

BDA Project: Hurricane forecasting in Stan

José Miguel Ramírez & Jonas Lindblad

Contents

Introduction	1
Modelling	3

```
source("init.r", local = knitr::knit_global())
```

Introduction

Tropical cyclones (in the Atlantic known as *hurricanes*) are destructive storms that occur during the late summer and fall in the northern hemisphere's tropical region. The storms are usually classified by their wind intensity in the region known as the *eye wall*, i.e. the circular region around the storm center where the wind speeds are strongest. Highly intense hurricanes can cause extreme levels of flooding in coastal areas and destroy buildings and homes. Furthermore, the monetary damages and loss of lives seem to increase with an almost exponential character as a function of storm intensity. As such, good probabilistic forecasts of hurricane intensity are important for purposes of evacuating areas at risk, as well as an interesting scientific problem.

Meteorologists forecast hurricane intensity changes with two classes of models: *dynamical* models, or *statistical* models. The dynamical models are based on systems of differential equations that are solved numerically with supercomputer clusters, while the statistical models are based on black-box machine learning methods or more basic statistical models. Rather bafflingly, one of the most successful models is a multiple linear regression model called Statistical Hurricane Intensity Prediction Scheme (SHIPS). Information about the SHIPS model and its datasets are available at the website: SHIPS Development.

The reason that a model as simple as SHIPS is still in operational use, is because of ongoing scientific challenges in understanding the complex causes of (as well as impediments to) intensification. It is well-known in the field that high sea temperatures in the tropical regions lead to dissipation and convection, which intensifies hurricane winds, while strong wind shear and surface drag causes weakening of the storms. However, knowledge of these factors have not lead to a complete solution to the problem and to some extent the issues may stem from the inherent complexity and uncertainty in current remote sensing technology.

In this project, we develop models inspired by the SHIPS scheme and using the SHIPS developmental data. The idea is to develop a Bayesian, probabilistic, intensity forecast that could (in principle) be put into operational use as a competitor to currently active statistical hurricane models. Not being affiliated with a forecasting agency has its limits however, as it is important to note that the SHIPS model also uses information from the numerical weather forecasts run by the US government. We do not have access to the global forecast models, and as such are limited to creating a *synoptic* model, i.e. an intensity forecast based purely on information available at the current time and not using data from supercomputers that simulate the laws of physics.

Data

This project uses the ‘Developmental Data’ from the SHIPS website: SHIPS Developmental Data. To be precise, we limit ourselves to the Atlantic version of the SHIPS data, and use the data from the 5-day SHIPS predictor file for the Atlantic. This data contains ~140 variables in 6-hourly time steps for every hurricane season 1982-2019.

Description of the variables

Describing all of the variables in the dataset is far outside the scope of this document, although a few important quantities should be known to the reader. Moreover, most of the variables are quite basic statistics computed by taking averages over certain regions (usually a circle or ring centered around the eye of the storm). The raw data used to compute the variables is a global grid discretization of the atmosphere, while the gridded discretization is created by statistically interpolating data from satellite remote sensing, weather balloons, weather stations, commercial/military aircraft, as well as weather buoys in the sea.

Include a picture example to go along with how the variables are computed?

Maximum wind speed (VMAX): This is the current maximum wind speed (in knots) in the eye wall of the hurricane, at ~100 meters above surface. Actually, there may be gusts and turbulences that are short-lived winds surpassing this wind speed, but it has been found that the maximum ~10 minute sustained winds in the eye wall is simpler to measure and predicts maximum gusts and storm damages best. In the Atlantic region this variable is measured by US military aircraft that drop sondes inside the eye wall of the storm. In case there is no scheduled flight mission to a hurricane, the maximum wind speed is instead inferred using statistics and satellite observations. Naturally, this is an estimate in either case, but the standard error is in fact significantly lower when dropsonde measurements are used. In all official US government data this quantity is rounded to the nearest multiple of five (this includes the SHIPS dataset). The DELTA12 variable corresponds to the value we want to predict: how the maximum wind speed will change in 12 hours.

Wind shear (SHRD): Wind shear is the phenomenon where the wind direction and speed differs significantly depending on height. The SHIPS dataset contains several wind shear variables, since wind shear is known to be an important variable that can impact the development of a tropical storm. In regions with high wind shear, it can be very difficult for a storm to intensify since excess energy is required for the storm to maintain stability. Sometimes high wind shear can displace and ‘rip apart’ a storm, causing it to weaken and potentially dissipate. Wind shear quantities are averages computed from wind data in the grid discretization of the atmosphere.

Sea temperature (CSST): The main energy source for tropical storms is the thermal energy in the sea, both the sea surface temperature and the deeper ocean heat content. The SHIPS dataset contains several variables representing the sea temperature, and these are well-known to correspond to storm intensification. The main physical process is that hot water evaporates, causing upward winds (convection) that lower the pressure inside the eye of the storm. This in turn causes winds to intensify as larger masses of the atmosphere get pushed towards the low-pressure point in the storm eye. These quantities are based on direct measurements from buoys and satellite remote sensing.

We will be studying the effect of different subsets of the data. The first data set will only contain the wind shear (SHRD) and the sea temperature (CSST), we will call from now on to this dataset as “B”. We will also study a dataset containing the same variables as in “A” plus an extra variable: (VMPI), we will be referring to this dataset as “B”. Finally, we will also test a larger dataset with the variables SHRD, CSST, VMPI, RHLO, and T200, we will refer to this dataset as “C”.

Prior work with the data

Although the data is publicly available with no restrictions, not much work using it can be found. This may be because of the lack of general interest in meteorology or possibly the poor state of the documentation. The authors of this report are not aware of any other work with this data apart from scientific publications by the American researchers involved in the SHIPS model development. The bulk of this research is available at the SHIPS website: SHIPS Model References.




Figure 1: Correlation plot of the variables we will include in the models.

Modelling

In this project we studied three different model structures with different variable selections. Drawing inspiration from the US government SHIPS model, all of our models are some variation of linear regression.

Description of the models

Linear regression

The first model is the simplest possible case - just a linear regression model. The mathematical statement of the model is as follows, where we let N denote the number of regression parameters.

$$y_i \sim \mathcal{N}(\alpha + X_i \cdot \beta_{N-1}, \sigma), \quad i = 1, \dots, r,$$

where we let X_i denote the i :th row of the data, β_{N-1} is an $N - 1$ -dimensional parameter vector, and r is the number of observations (rows) in the data. The inference is done on the parameters α, β, σ . They are given the following (weak) prior distributions.

$$\begin{bmatrix} \alpha_0 \\ \beta_{N-1,0} \end{bmatrix} \sim \mathcal{N}(\mathbf{0}_N, 10 \cdot \mathbf{I}_N), \sigma_0 \sim \text{Inv-}\chi^2(\frac{1}{10}).$$

The Stan code for this model is the following one:

```
writeLines(readLines("models/linear.stan"))
```

```
##
## data {
##   int<lower=0> N;
##   int<lower=0> J;
##   vector[N] y;
##   matrix[N,J] x;
##   vector[J+1] mu; // required prior means
##   matrix[J+1, J+1] tau; // prior covariance matrix
```

```

## }
##
## parameters {
##   vector[J+1] theta;
##   real < lower =0 > sigma;
## }
##
## model {
##   theta ~ multi_normal(mu, tau);
##   sigma ~ inv_chi_square(0.1);
##   y ~ normal( theta[1] + x*theta[2:J+1], sigma);
## }
##
## generated quantities {
##   vector[N] log_lik;
##   for (n in 1:N) {
##     log_lik[n] = normal_lpdf(y[n] | theta[1] + x[n]*theta[2:J+1], sigma);
##   }
}

```

We run this model with 4 chains for 4000 iterations using 2000 as warm-up. We used the datasets “A”, “B”, and “C”, and for all of them, the ESS is always higher than 100, and the \hat{R} values approximated to 1.0 for each parameter, which means that we run enough iterations to reach convergence. If this was not the case and the \hat{R} ratios were higher than 1, the chains would not be sampling from the same part of the parameter space and we would need either to increase the number of iterations or to edit the proposal distribution. The definition of \hat{R} that we used can be found here.

Skew-regression

The second model is a variation of the first, where instead of a normal distribution around the regression line, a skew-normal distribution is used. This probability distribution has different properties, making computation a bit less efficient, but there are scientific reasons to believe that a skew towards higher intensity changes would be a better fit. The model is summarized in mathematical language as follows.

$$y_i \sim \text{SkewNormal}(\alpha + X_i \cdot \beta_{N-1}, \sigma, \psi), \quad i = 1, \dots, r,$$

where we use the same conventions as in the linear model. The priors of the parameters are given below.

$$\begin{bmatrix} \alpha_0 \\ \beta_{N-1,0} \end{bmatrix} \sim \mathcal{N}(\mathbf{0}_N, 10 \cdot \mathbf{I}_N), \sigma_0 \sim \text{Inv-}\chi^2\left(\frac{1}{10}\right), \psi_0 \sim \mathcal{N}(0, 1).$$

We run this model for the same number of iterations as before: 4000 using with 2000 as warm-up, the same number of chains: 4, and the same datasets: “A”, “B”, and “C”. But in this case we also increased the maximum tree depth to 15. Convergence was also reached in all these cases, as all \hat{R} values approximated to 1.0 and all ESS were higher than 100.

Regression with changing variance

This is another variation of the linear regression model. In this model there is also a regression variance dependence on the wind intensity V_{max} . This somewhat justified, since very extreme values of V_{max} lead to volatile situations where the evolution of the storm can be difficult to predict even for the best forecasting agencies. There are also relatively few cases of high V_{max} since every storm starts out weak, but only a few evolve to be high-category hurricanes.

This model is declared in mathematical notation as follows:

$$y_i \sim \mathcal{N}(\alpha + X_i \cdot \beta_{N-1}, \sigma + \gamma|V_{max,i}|), \quad i = 1, \dots, r,$$

where, again, we use the same notation as before and let $V_{max,i}$ denote the V_{max} -value of the i :th row. Note the absolute value function in the variance section; this statement is used to ensure that the variance remains positive. The parameters are given the priors listed below.

$$\begin{bmatrix} \alpha_0 \\ \beta_{N-1,0} \end{bmatrix} \sim \mathcal{N}(\mathbf{0}_N, 10 \cdot \mathbf{I}_N), \quad \sigma_0 \sim \text{Inv-}\chi^2\left(\frac{1}{10}\right), \quad \gamma_0 \sim \Gamma(1, 1).$$

To execute this model we used 4 chains, 4000 iterations with half of them as warm-up and a maximum tree depth of 15. We checked the \hat{R} values for the three variable combinations (“A”, “B”, and “C”), and all of them approximated to 1.0, so we can state that the chains converged in the three cases. Moreover ESS is always high enough.

Models Comparison

We assessed the models using leave-one-out cross-validation (LOO-CV) to do further comparisons, and we used the `r` package `loo` for using PSIS-LOO as method to approximate the exact LOO given the posterior samples.

The PSIS-LOO elpd for the basic linear model using data “A” is -1636.0 with 9 as effective number of parameters (`p_loo`), looic of 3272.1 and Monte Carlo Standard Error (MCSE) of 0.0. For data “B”, the PSIS-LOO elpd is -1633.6 with `p_loo` = 10, looic = 3267.2 and MCSE = 0.1. For the last dataset “C”, the elpd_loo is -1608.6, with `p_loo` = 13.1, looic = 3217.2 and MCSE = 0.1. All the Pareto k estimates are lower than 0.5 for the first two datasets, and there is only one value for “C” above 0.5, but it is lower than 0.7, as we can see in Figure. X. This means that the three PSIS-LOO estimates can be considered to be reliable.

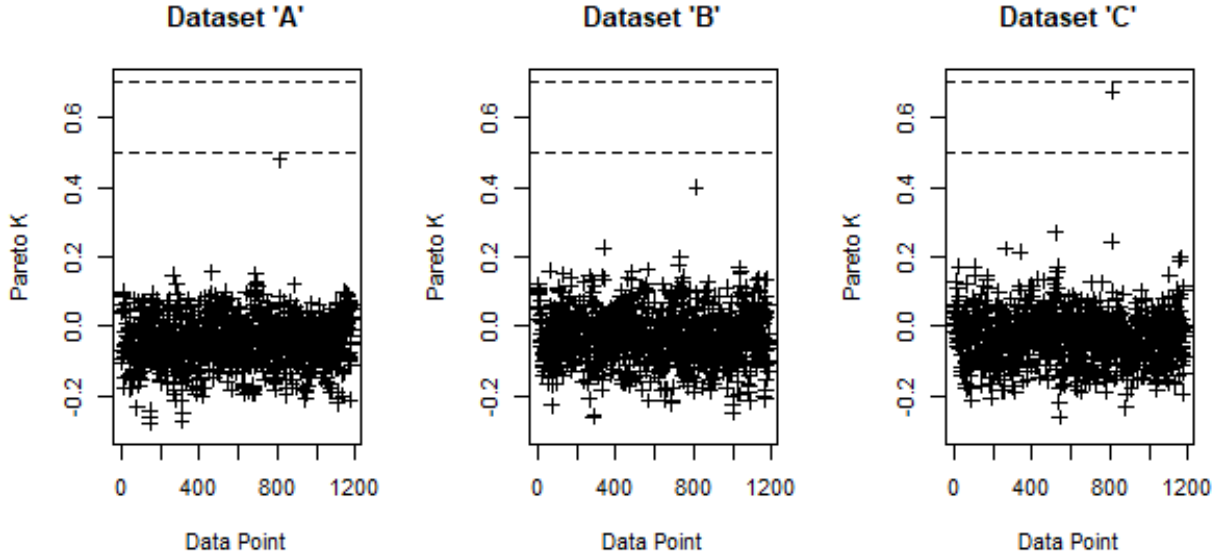


Figure 2: PSIS-LOO diagnosis for the linear model.

The model with dataset “C” has the highest PSIS-LOO and the lowest looic associated to it, but in order to better assess the differences between the datasets we have used the function `\textit{loo_compare}`. This function computes the difference in elpd relative to the model with the largest elpd (in this case, the “C” model). As we can see in Table X, the differences in elpd and the scale relative to the standard error of these differences indicate that the best dataset for this model according to PSIS-LOO would be the dataset “C”.

variable set	elpd_diff	se_diff
<i>C</i>	0.0	0.0
<i>B</i>	-25.0	6.5
<i>A</i>	-27.4	6.3

The PSIS-LOO elpd for the skewed-regression model using data “A” is -1608.1 with $p_{\text{loo}} = 9.6$, $\text{looic} = 3272.1$, and $\text{MCSE} = 0.1$. For dataset “B”, the PSIS-LOO elpd is -1602.5 with $p_{\text{loo}} = 10.4$, $\text{looic} = 3205.1$ and $\text{MCSE} = 0.1$. For the last dataset “C”, the elpd_loo is -1579.4, with $p_{\text{loo}} = 13.3$, $\text{looic} = 3158.7$ and $\text{MCSE} = 0.1$. All the Pareto k estimates are lower than 0.5 for the first two datasets, and there is only one value for “C” above 0.5, but it is lower than 0.7, as we can see in Figure. X. This means that the three PSIS-LOO estimates can be considered to be reliable.




Figure 3: PSIS-LOO diagnosis for the skewed model.

Using `\textit{loo_compare}` we can observe that the best dataset for this model would be again dataset “C” (Table X).

variable set	elpd_diff	se_diff
<i>C</i>	0.0	0.0
<i>B</i>	-23.2	6.2
<i>A</i>	-28.7	6.2

The PSIS-LOO elpd for the regression with changing variance using data “A” is -1440.2 with $p_{\text{loo}} = 7.3$, $\text{looic} = 2880.5$, and $\text{MCSE} = 0.1$. For dataset “B”, the PSIS-LOO elpd is -1435.8 with $p_{\text{loo}} = 8.3$, $\text{looic} = 2871.5$ and MCSE approximated to 0.0. For the last dataset “C”, the elpd_loo is -1403.1, with $p_{\text{loo}} = 12.0$, $\text{looic} = 2806.2$ and $\text{MCSE} = 0.0$. All the Pareto k estimates are lower than 0.5, as we can see in Figure. X, which means that the three PSIS-LOO estimates can be considered to be reliable.

Using again the function `\textit{loo_compare}` we can see that the best dataset for this model would be again dataset “C” (Table X).




Figure 4: PSIS-LOO diagnosis for the variance model.

variable set	elpd_diff	se_diff
<i>C</i>	0.0	0.0
<i>B</i>	-32.6	8.2
<i>A</i>	-37.1	8.2

We have seen that the best dataset for all of our Stan models is dataset “C”. If we compare now the three models using this dataset, we can see that the best performing model according to PSIS-LOO would be the last studied model, the regression with changing variance.

Model	elpd_diff	se_diff
Variance	0.0	0.0
Skew	-176.3	27.8
Linear	-205.5	34.9